

STAT:5400 Section 1.1 — Introduction

Boxiang Wang

2026-08-26

1 Why statistics & data science?

- 2009: *“I keep saying that the sexy job in the next 10 years will be statisticians. And I’m not kidding.”*

[Hal Varian](#), Chief Economist at Google

- 2012: [“Data Scientist: The Sexiest Job of the 21st Century”](#), Harvard Business Review.
- 2022: [“Is Data Scientist Still the Sexiest Job of the 21st Century?”](#), Harvard Business Review.
 - The U.S. Bureau of Labor Statistics predicts data science will see more growth than almost any other field between now and 2029. Data scientist jobs are projected to grow **+36%** (2023–2033), with over 20,000 openings per year.
 - In 2024, the [median pay was \\$112,590](#). The median salary for an experienced data scientist in California is approaching \$200,000.
 - Data Scientist is [#6 Best STEM Job](#) and [#4 Best Technology Job](#).
- 2023–2025: the question changes. Nobody has written the next “sexiest job” piece; the discussion is now whether AI changes the job. [A 2023 essay](#) argued the title is “being challenged by AI”; [a 2025 essay](#) called it “the rise, fall, and evolution” of the job.
- 2026: the numbers say the job kept growing. [BLS projects](#) data-scientist employment **+34%** over 2024–2034, about 23,400 openings a year; [for statisticians](#), growth “much faster than average” as well.

2 Data Science

- The science of data science focuses on the computation and inference for the **data-driven analysis** from the interdisciplinary perspectives.
- The term **data science** is believed to be credited to [Jeff Wu](#), Georgia Tech.
- Statistical endeavors:
 - Applied statistics and data analysis
 - Development of statistical methods and software
 - Research in statistical theory
- Computing essential to all of statistical endeavors.
 - Computational intensive methods.
 - Efficiency facilitate modern statistical research.
 - Large-scale data pipelines; GPU/TPU-accelerated algorithms.


3 Data science: skills and jobs

- **Core skill set.**
 - *Statistics*: modeling, inference, experimental design, uncertainty quantification
 - *Computing*: Python/R, data engineering, GPU-accelerated methods, versioning
 - *Communication*: domain knowledge, visualization, storytelling, ethics

What employers ask for (checked Aug 2026):

- **Statistician** — [O*NET, U.S. Department of Labor, 15-2041](#): top skills *mathematics, critical thinking, reading, listening, speaking*; tools employers list: **Python, R**, SAS, SPSS, Tableau, Excel.
- **Data scientist** — [O*NET, 15-2051](#): top skills *mathematics, critical thinking, active learning, speaking*; tools: **Python, SQL**, Git, Spark, TensorFlow/PyTorch, AWS, Tableau.
- **Data-scientist job postings** — [827 U.S. postings, 2025–26](#): machine learning 69%, **Python 57%, R 33%, SQL 30%**, NLP 19% (5% a year earlier); a graduate degree in 54% (PhD 24%, master's 30%).
- **What is rising** — [LinkedIn Skills on the Rise 2026](#): AI literacy, data-driven decision making, data storytelling, data governance and responsible AI. [Indeed](#): 45% of data & analytics postings now mention AI.
- **What people use** — [Stack Overflow Developer Survey 2025](#): Python 58%, SQL 59%, R 5%; 84% use or plan to use AI tools, 33% trust the answers.
- Live listings: [Glassdoor, data scientist](#) · [BLS, data scientists](#) · [BLS, statisticians](#).

Read the two O*NET lists again: both start with **mathematics and critical thinking**; the tools come after. Your degree is the first part. This course is the second: **Python, R, SQL, Git**, and the habit of checking.



3.1 A short history of generative AI

- 2017: the transformer.
- 2018–2020: GPT-1, BERT, GPT-2, then GPT-3 with 175 billion parameters. Trained on text; it saw little code.
- 2021: [Codex](#) and GitHub Copilot. Adding code to training data made models better at reasoning, not just at code. That is one reason “AI can code” arrived before “AI can reason”.
- November 30, 2022: ChatGPT. 100 million users in two months.
- 2023: GPT-4 (March), Llama with open weights, Claude, Gemini (December).
- 2024: GPT-4o, Claude 3, Gemini 1.5 with a million-token context, the first “reasoning” models (o1, September). Also the year [ChatGPT said 9.11 is larger than 9.9](#).
- 2025: DeepSeek-R1 (January): open weights, a fraction of the cost. Agents and tool use everywhere. GPT-5 (August).
- 2026: GPT-5.6, Claude 5, Gemini 3.6, DeepSeek V4, and hundreds more. The model you use this semester will be replaced before you graduate; the statistics will not.

The short version, for this course. For the long one, IBM has a good [history of AI](#).

4 AI for statistics, statistics for AI

Two directions, and this course runs both ways.

- **AI for statistics** — the model as a tool in the statistician’s workflow.
 - Writing and translating code: R → Python, a loop → a vectorized version, a formula → a function (Section 2).
 - Cleaning and extracting: messy text → a tidy table; structured output instead of prose (Section 6).
 - Designing and explaining: draft a simulation study, explain an error message, summarize a paper — then *you* verify.
 - Learning: the *Ask AI* buttons in these notes, and the course tutor.
 - Agents: an AI that runs a small analysis end to end, from data to report, while you watch each step. In class we replay [one on MorphMind](#) (free tier, no payment required).

Statistical agents you can try today:

- [ChatGPT](#): upload a file, it writes and runs Python and shows you the code.
- [Claude](#): code execution in the chat, or Claude Code in a terminal.
- [Gemini’s Data Science Agent in Colab](#): describe the goal, get a whole notebook.
- [MorphMind AgentLab](#): a team of agents with receipts for every step.
- Open source: [MetaGPT’s Data Interpreter](#), [PandasAI](#), and agent frameworks such as LangGraph, CrewAI, and the OpenAI Agents SDK.
- From R: [e1lmer](#) with tool calling (Section 6) lets you build a small one yourself.
- **Statistics for AI** — the model as an object of statistical study.
 - An LLM is a sampler: temperature and top-p are the same ideas as random-variate generation (Section 3).
 - Evaluating a model is a simulation study: accuracy is an estimate, needs a confidence interval, and “how many test prompts?” is a sample-size question (Section 3).
 - Bootstrap a benchmark score; test whether model A really beats model B (Section 4).
 - Calibration, hallucination rates, and uncertainty for a black box (conformal prediction) — the statistician’s contribution to trustworthy AI (Section 6).
- **Statistics has an edge here.** An LLM is a probability model: it generates text from a distribution and predicts the next word. Generating from a distribution, prediction, inference: these are the questions statistics has worked on for a century.
- **Then why are statisticians not playing a larger role in AI?** Not for lack of ideas. What is missing is often the skills: large-scale computation, modern coding practice, text as data. This course is about closing that gap.

Reference: *An Overview of Large Language Models for Statisticians*, The American Statistician (2025), arXiv:2502.17814.

5 Why learn to code when AI can code?

- **Interviews have no AI in them.**
 - [Anthropic bans it in live interviews](#) and has dropped candidates who used it.
 - [Google is moving interviews back in person](#).
 - The by-hand blocks and the midterm here are the same test.
- **Many employers do not allow it.**
 - [The big banks](#), [Samsung](#), and Apple restricted ChatGPT.
 - Patient records, clinical trials, government work: the data cannot leave the building.
 - [At Iowa](#): university data goes only into Copilot Chat with your HawkID, and a person checks every output.
- **AI code is often almost right. Almost is your job.**
 - [66% of developers](#) say their biggest problem is code that is almost right.
 - [A randomized trial](#) found experienced developers 19% slower with AI, while they believed they were faster.
 - [AI-written code picked the insecure option 45% of the time](#).
 - A wrong variance formula or an unset seed still runs. Only someone who can read the code will catch it.
- **Big models are expensive.**
 - [Uber's engineers](#) ran up \$500 to \$2,000 a month each on AI coding tools.
 - [Klarna](#) replaced 700 support staff with AI, then hired people back.
 - By 2026 the headline was "[AI can cost more than the human workers](#)".
 - [NVIDIA](#) argues agents should run on small, specialized models. Someone has to build them. That is programming plus statistics.
- **You cannot supervise what you do not understand.**
 - The job is moving from writing code to checking it.
 - [Deloitte refunded the Australian government](#) for a report with invented citations.
 - [Courts keep fining lawyers](#) for citing cases that do not exist.
 - Statistics is the discipline of knowing when to trust an estimate. An AI answer is an estimate.

6 Goals of this course are to develop

- Intelligent use of appropriate computing tools.
 - R and Python; SAS, C, Fortran when needed.
 - AI as a tool: chatbots, APIs, agents, and knowing when not to trust them.
- Use of software engineering techniques for scaling projects.
 - Version control: git and GitHub.
 - Parallel computing.
- Understanding of important statistical computing algorithms.
 - Bisection search, Newton's method, gradient descent.
 - Bootstrap, cross-validation.
- Experience of computational intensive methods.
 - High-dimensional statistics.
 - Machine learning methods.
- Ability to design and implement simulation studies.
- Communication of statistical ideas.
 - LaTeX, R Markdown, Quarto.
 - Writing R packages.

7 Set up your accounts

► Run — first, click *Run Code*. This runs in your browser; nothing to install.

```
x <- rnorm(100)
mean(x)
```

Do these this week; the notes assume them from Section 2 on.

1. **GitHub account** — <https://github.com/signup>. Add and verify your @uiowa.edu email under Settings → Emails.
2. **GitHub Education** (free Copilot and Pro-level Codespaces) — <https://github.com/settings/education/benefits> → *Start an application*. Proof: a current student ID, class schedule, or enrollment letter; using the @uiowa.edu email speeds review. When approved, GitHub Copilot Student and Codespaces are activated on your account.

3. **Test the Codespace** — click *Open in Codespace* at the top of this page. The first start builds the course environment (a few minutes); later starts take about 30 seconds. Open `notes/run/s1p1_sumC.R`, put the cursor on line 3, press Cmd/Ctrl+Enter.
4. **One chatbot login** — ChatGPT, Claude, or Gemini; the free tier is enough. The *Ask AI* buttons in these notes open one of them with the question filled in.
5. **A Gemini API key** (free) — <https://aistudio.google.com/apikey>. Used by the first API call in Section 17. Details in the box below.
6. **OpenAI and Anthropic keys** (optional, paid; \$5 is enough for the course). Details below.
7. **University systems** — the department's Linux servers `l-lnx2xy.stat.uiowa.edu` (HawkID, ssh; VPN from off campus; Tech Guide in Section 2); [IDAS](#) for the midterm; [Argon](#), the campus cluster, if you want it (I request class access; for research ask your advisor).
8. **MorphMind account** (free) — <https://morphmind.ai>. The course tutor and the agent demos run there.

Never paste an API key into a script, a notebook, or these notes. A key in a public repository is found and abused within minutes.

i Gemini API key, step by step (free)

1. Go to <https://aistudio.google.com/apikey> and sign in with a Google account (your personal one is fine).
2. Click **Create API key**. If it asks for a project, accept the default. Copy the key; it looks like `AIza...`
3. Store it on GitHub: profile picture → **Settings** → **Codespaces** → **New secret**. Name `GEMINI_API_KEY`, paste the key as the value, under **Repository access** pick the course repo, **Add secret**.
4. Restart your Codespace (it only reads secrets at start). Check in R: `Sys.getenv("GEMINI_API_KEY")` should not be empty.
5. Free tier: a few requests per minute and a daily cap, plenty for this course. If you hit the limit, wait a minute.
6. HTTP 503 **Service Unavailable** means Google's servers are busy, not that your key is wrong. Retry, or pick another model: `models_google_gemini()` lists what your key can use; `chat_google_gemini(model = "gemini-3.7-flash")` chooses one. The Ollama models below never have this problem.

i OpenAI API key (paid)

1. <https://platform.openai.com> → sign up (this is separate from a ChatGPT subscription).
2. **Settings** → **Billing** → add a payment method and buy \$5 of credit. Nothing runs without credit.
3. **API keys** → **Create new secret key** → copy it now; it is shown once.
4. Codespaces secret named `OPENAI_API_KEY`, same steps as above. Restart the Codespace.
5. In R: `library(ellmer); chat_openai()$chat("Say hi.")`. A one-line answer costs a fraction of a cent.

i Anthropic (Claude) API key (paid)

1. <https://console.anthropic.com> → sign up (separate from a Claude.ai subscription).
2. **Billing** → buy \$5 of credit.
3. **API Keys** → **Create Key** → copy it now.
4. Codespaces secret named `ANTHROPIC_API_KEY`. Restart the Codespace.
5. In R: `library(ellmer); chat_anthropic()$chat("Say hi.")`.

i Small models with no key: Ollama

No key, no cost, works offline. The course Codespace already has Ollama with two models: `qwen2.5:0.5b` (Alibaba, 0.5 billion parameters, 0.4 GB) and `llama3.2:1b` (Meta, 1 billion, 1.3 GB). On your own computer, Ollama is one download.

1. On your laptop: install from <https://ollama.com> (Mac, Windows, Linux), then `ollama pull qwen2.5:0.5b` and `ollama pull llama3.2:1b` in a terminal. In the Codespace, skip this.
2. From R, `ellmer` talks to a local model the same way it talks to Gemini. Only the function name changes; Ollama is a small server on your own machine.

```
library(ellmer)
small <- chat_ollama(model = "qwen2.5:0.5b")
bigger <- chat_ollama(model = "llama3.2:1b")
small$chat("Which is larger, 9.11 or 9.9? Answer in one line.")
bigger$chat("Which is larger, 9.11 or 9.9? Answer in one line.")
```

Ask the small one five times. It still gets this wrong some of the time; that is the 2024 story in Section 3, on a machine you control. The 1B model is slower and

more often right. Bigger models need more memory than a Codespace has; on your laptop try llama3.2 (3B).

3. From Python:

```
import ollama; ollama.chat(model="qwen2.5:0.5b", messages=[{"role": "user", "content": "..."}])
```

.
4. Also worth knowing: the R package `localLLM` loads a model inside R with no server at all. We will use it in the Tech Guide.

7.1 Where does code run?

- **On this page.** Most examples run in your browser. Click *Run Code*.
- **In a Codespace.** Examples that need a compiler, an API key, or more memory have a *Run in Codespace* button. Same code, real machine.

7.2 Optional: install on your own laptop

Not required for this course. If you want a local setup, use this one so we all match. Each item has a folded box with the steps, and an *Ask AI* button whose prompt asks for today's instructions for your machine.

i R and RStudio (or Positron)

1. R: <https://cran.r-project.org> → your operating system → the latest release. Mac: pick the Apple silicon build if your Mac is M1 or newer.
2. RStudio: <https://posit.co/download/rstudio-desktop/>. Or Positron, Posit's newer editor for R and Python: <https://positron.posit.co>.
3. Check: open RStudio, run `R.version.string`.
4. Packages we use early: `install.packages(c("microbenchmark", "Rcpp", "ellmer"))`. Rcpp needs a compiler: Mac, `xcode-select --install`; Windows, Rtools from CRAN.

i Python with uv

1. Install uv: <https://docs.astral.sh/uv/getting-started/installation/> (one command for Mac, Linux, or Windows).
2. `uv python install 3.12`
3. In a project folder: `uv init`, then `uv add numpy pandas matplotlib jupyter google-genai`.
4. Check: `uv run python -c "import numpy; print(numpy.__version__)"`.

i VS Code, Quarto, and Copilot

1. VS Code: <https://code.visualstudio.com>.
2. Extensions (View → Extensions): Quarto, R, Python, Jupyter, GitHub Copilot.
3. Quarto: <https://quarto.org/docs/get-started/>. Check: `quarto check` in a terminal.
4. Sign in to GitHub inside VS Code so Copilot works (free with GitHub Education).


i git

1. Mac: `xcode-select --install` in Terminal gives you git and a C compiler. Windows: <https://git-scm.com/download/win>. Linux: `sudo apt install git`.
2. Tell git who you are: `git config --global user.name "Your Name"` and `git config --global user.email "you@uiowa.edu"`.
3. Check: `git --version`.

8 Applied Statistics

Application of some statistical tools to analyze a practical problem.

Data collection → modeling → decision making.



```
> library(faraway)
> glm(cbind(damage, 6 - damage) ~ temp,
+     family=binomial, oranges)


Call: glm(formula = cbind(damage, 6 - damage) ~ temp, family = binomial,
data = oranges)

Coefficients:
(Intercept)      temp 
 11.6630      -0.2162 

Degrees of Freedom: 22 Total (i.e. Null);  21 Residual
Null Deviance:      38.9 
Residual Deviance: 16.91   AIC: 33.67
```

9 Optimization


Translate a conceptual problem into a computational problem.



Logistic regression:

$$P : \min_{\beta} \left[- \sum_{i=1}^n y_i \mathbf{x}_i' \beta + \sum_{i=1}^n \log(1 + e^{\mathbf{x}_i' \beta}) \right].$$

1. In this class, we may learn to call some general routines to solve Problem P directly.
2. We also want to develop and implement some algorithms specifically for solving Problem P , in order to obtain
 - knowledge of the performance of different algorithms.
 - deeper understanding of the statistical problem.



An application. On January 28, 1986, the space shuttle Challenger broke apart 73 seconds after launch. The night before, engineers argued about the O-ring seals and the forecast of 31°F. Here are the 23 earlier launches: temperature at launch and how many of the six O-rings showed damage. Problem P with one predictor.

► Run


```

      c(median_ms = 1000 * median(e1), min_ms = 1000 * min(e1))
    )))
rownames(out) <- sapply(exprs, function(e) paste(deparse(e), collapse = ""))
round(out, 2)
}

```

► Run

```

x <- runif(1e6)
timeit(
  sqrt(x),
  x ^ 0.5
)

```

✎ Edit — change the vector length or add your own pair of expressions, then rerun.

```

x <- runif(1e6)
timeit(
  x * x * x,
  x ^ 3,
  x * 0.1,
  x / 10
)

x <- matrix(2, 1000, 1000)
timeit(
  colMeans(x),
  rowMeans(x)
)

```

► Run — the same product, two ways to put the parentheses.

```

A1 <- matrix(2, 1000, 1000)
A2 <- matrix(2, 1000, 1000)
A3 <- rep(2, 1000)

system.time(A1 %*% A2 %*% A3)
system.time(A1 %*% (A2 %*% A3))

```

One more: `mean(x)` vs `sum(x) / length(x)`. Two ways to compute an average. Guess which is faster, then run it.

```

#| autorun: true
#| edit: false
import numpy as np, timeit
def tm(times=5, **fns):
    """tm(): the Python twin of timeit(); median wall-clock time in ms"""
    for name, f in fns.items():
        t = timeit.repeat(f, number=1, repeat=times)
        print(f"{name:>16s}: {1000*np.median(t):9.2f} ms")

```

 Edit

10.1 R

```

x <- runif(1e6)
timeit(mean(x), sum(x) / length(x))

```

10.2 Python

```

x = np.random.rand(1_000_000)
tm(mean=lambda: x.mean(), sum_over_n=lambda: x.sum() / len(x))

```

`mean()` is about twice as slow. It is not lazy; it makes two passes: a first average, then a correction using the deviations from it. The slower one is the more accurate one. Add up a million copies of 0.1 and see:

► Run

```

x <- rep(0.1, 1e6)
s <- 0
for (v in x)
  s <- s + v
c(loop = s / length(x) - 0.1, sum_over_n = sum(x) / length(x) - 0.1, mean = mean(x) - 0.1)

```

The loop adds in plain double precision and drifts by about $1e-12$; on some machines `sum()` does too. `mean()` comes back with exactly 0. More on floating point in Section 2; more speed tricks as we need them.

Run in Codespace — a plain R loop vs C++ (through Rcpp) vs the built-in `sum()`. A browser cannot compile C++.

```
library(Rcpp)
library(microbenchmark)
sumR <- function(x) {
  total <- 0
  for (i in seq_along(x))
    total <- total + x[i]
  total
}
cppFunction('double sumC(NumericVector x){
  int n = x.size();
  double total = 0;
  for (int i=0; i < n; i++) {
    total += x[i];
  }
  return total;
}')
x = runif(1e3)
microbenchmark(sum(x), sumC(x), sumR(x))
```

Source: *Advanced R*.

👉 By hand — write `sumR()` yourself (a `for` loop that adds up the elements of `x`), then time it against the built-in `sum()`. Done when `all.equal()` returns `TRUE` and the timing table appears. No AI for this block.

```
sumR <- function(x) {
  # your code here
}

x <- runif(1e5)
all.equal(sumR(x), sum(x))
timeit(sum(x), sumR(x))
```

Solution

```
sumR <- function(x) {
  total <- 0
  for (i in seq_along(x)) total <- total + x[i]
  total
}
```

🐙 Vibe — AI-assisted extension: turn the timing into a function of `n` and plot how the $x^{0.5}$ vs `sqrt(x)` ratio changes with `n`. Starter file for Copilot: [notes/vibe/slp1_benchmark.R](#) (open it in the Codespace). Prefer no AI? Write it here.

```
bench_pow <- function(n) {  
  # return: median time of x^0.5  
  #           / median time of sqrt(x)  
  
}  
# n_grid <- 10^(4:6); ratios <- sapply(n_grid, bench_pow)  
# plot(n_grid, ratios, log = "x", type = "b")
```

11 Types of computer products

- **Operating systems: make the computer itself work.**
e.g. Linux, Windows, Unix, MacOS...
- **Applications: perform specific tasks.**
e.g. Microsoft Word, Excel, S-Plus, R, SAS ...
- **Commercial Systems.**
 - code and lists of bugs are secret
 - expensive
 - require upgrading and relicensing
 - Microsoft products, S-Plus, SAS, SPSS, Unix, etc.
- **Free Open Source systems.**
 - revolution in software availability and function from the open source movement
 - can see all code, change it, learn from it
 - quality generally quite good
 - more rapid updates
 - most products have an active and helpful user news group
 - generally lack some fancy features like extensive GUI
 - Linux, LaTeX, R

Source: Frank Harrell (University of Virginia) lecture note.

12 User interfaces: graphical vs command line


- **graphical (GUI, mouse, menus)**
 - easier to learn

- less flexible
- repetitive when the same tasks have to be repeated
- hard to document the exact steps taken
- hard to reproduce results


- **command line interfaces**

- harder to learn
- more flexible and powerful
- can save commands in scripts to replay when the same tasks have to be performed repeatedly
- can write generic commands to facilitate running different analyses with the same structure


13 GUI vs command line: text editor: sublime vs vim




14 GUI vs command line: git (version control)



15 GUI vs command line: remote access to servers



16 R vs Python




Where each language lives (checked Aug 2026):

- **Python:** the AI ecosystem.
 - #1 on [TIOBE](#).
 - Used by 58% of developers in the Stack Overflow 2025 survey.
- **R:** statistics and research.
 - #9 on [TIOBE](#).
 - Taught in 20 of 27 design-of-experiments courses surveyed at U.S. universities.
 - Articles citing R: about 1,000 in 2005, 62,000 in 2019.
 - Exception: economics and business journals still run on Stata.
- **SAS:** regulated pharma.
 - FDA has said since 2015 that it [does not require any specific software](#). SAS stayed the default by habit.
 - [Novo Nordisk](#) and [Roche](#) have filed R-based submissions; the [R Consortium's FDA pilots](#) have all been accepted.
 - SAS is about to leave TIOBE's top 30 for the first time. Knowing both is an edge in that industry.
- **Rust:** the engine under the tools.
 - You will meet it without writing it: [Polars](#), [uv](#), and [ruff](#) are Rust inside, Python outside, the same way R uses C++ through Rcpp.

- #10 on TIOBE; the most admired language on Stack Overflow ten years running.
- **Julia**: excellent for numerical work, real in scientific-computing groups.
 - Never a stable top-30 language; rarely in job postings. A specialty, not a foundation.

17 LLM: chatbox vs API



Three ways to use a model, in the order this course meets them:

- **Chatbox** — a person talks to the model. You already do this. Good for questions, explanations, drafts.
- **API** — a program talks to the model (Session 6). The answer comes back as data: you can loop over 1,000 prompts, set the temperature, ask for a table instead of prose.
- **Agent** — the model calls tools and runs steps on its own, and you supervise (Session 19). The skill is no longer writing the answer; it is specifying the task and checking the work.

17.1 Your first API call

The chatbox is a person talking to a model. The API is a *program* talking to a model: the same question, but the answer comes back as data your code can use. That is the door to everything in Section 6.

Run in Codespace — R, with the `ellmer` package (needs your key):

```
library(ellmer)
models_google_gemini() # which models your key can use; pick a "flash" one
chat <- chat_google_gemini(model = "gemini-3.7-flash") # reads GEMINI_API_KEY from the env.
chat$chat("In one sentence, what does a statistician do that a chatbot cannot?")
```

Run in Codespace — Python, with Google's `google-genai` package:

```
from google import genai
client = genai.Client() # reads GEMINI_API_KEY from the environment
interaction = client.interactions.create(
    model="gemini-3.7-flash",
    input="In one sentence, what does a statistician do that a chatbot cannot?"
)
print(interaction.output_text)
```

Three things to notice, all of which return in Section 6:

- The model is a function call with arguments — here the model name and the prompt; later temperature, system instructions, and structured output.
- Run it twice and you get two different answers. An LLM is a *sampler*; Section 3 is about samplers.
- The answer is text, not truth!

18 Before next class

- ☐ Create a GitHub account and verify your `@uiowa.edu` email.
- ☐ Apply for GitHub Education.
- ☐ Open the course Codespace once and run `notes/run/s1p1_sumC.R`.
- ☐ Create a Gemini API key and store it as the `GEMINI_API_KEY` Codespaces secret.
- ☐ Run the two API examples above in the Codespace. Did the two answers differ?

Questions about these notes? Later this semester the course tutor, built on [MorphMind](#), will answer them. Until then, the *Ask AI* buttons.